

Chakravat Manthan — Comprehensive Machine Learning Technical Report

Project Title: Chakravat Manthan (■■■■■■■■ ■■■■) — Operational Tropical Cyclone Intensity & Trajectory Intelligence Platform

Domain: Satellite Meteorology, Spatio-Temporal Deep Learning, Numerical Atmospheric Physics

Data Sources: ISRO MOSDAC INSAT-3DR Geostationary Satellites, IMD Best Track Archives, NOAA/IBTrACS

Report Date: September 27, 2026

Authors: Chakravat Manthan AI Engineering Team

Executive Summary

Tropical cyclones across the North Indian Ocean (Bay of Bengal and Arabian Sea) represent some of the deadliest meteorological events on Earth due to high coastal population density, shallow bathymetry, and complex land-sea interactions. Traditional operational forecasting has relied predominantly on manual subjective techniques (e.g., the Dvorak Technique) or compute-heavy Numerical Weather Prediction (NWP) models (e.g., WRF, GFS) which experience multi-hour latency.

The **Chakravat Manthan** project was initiated to build an automated, end-to-end, deep-learning-driven operational intelligence engine capable of:

1. Ingesting raw geostationary thermal infrared (TIR-1) satellite radiance from ISRO's INSAT-3DR payload.
2. Directly regressing continuous physical cyclone parameters (Maximum Sustained Wind Speed $V_{\max}$, Central Atmospheric Pressure $P_{\min}$, and Spatial Swath Dimensions).
3. Classifying cyclone intensity according to the official 8-stage India Meteorological Department (IMD) scale with calibrated softmax probability distributions.
4. Extrapolating track trajectory and forecasting inland penetration (including storm decay over interior states like Odisha, Jharkhand, and Chhattisgarh).

This document serves as the comprehensive whitepaper detailing the genesis, engineering roadblocks, mathematical breakthroughs, architecture design, and production readiness of the machine learning system.

1. Project Genesis & Starting Point

1.1 The Baseline Problem

Initial efforts in deep learning for tropical cyclones typically treat cyclone satellite imagery as isolated static 2D image classification tasks. However, this naive approach fails fundamentally in operational meteorology due to:

- **Memorylessness:** A storm with 65-knot winds undergoing rapid intensification looks drastically different in its convective evolution than a weakening 65-knot storm post-landfall. Single-frame 2D CNNs cannot disambiguate the temporal derivative of intensity ($\frac{dI}{dt}$).
- **Categorical Discretization Loss:** Forcing continuous atmospheric dynamics into discrete integer labels throws away critical gradient information between a borderline 63-knot Severe Cyclonic Storm (SCS) and a 64-knot Very Severe Cyclonic Storm (VSCS).
- **Sensor Noise & Satellite View Angle:** Geostationary satellites over the Indian Ocean view storms at non-zero zenith angles, causing limb-cooling artifacts and cloud parallax errors.

1.2 Initial Architecture & Objectives

We established the goal of creating a **Spatio-Temporal Deep Neural Network** operating directly on multi-frame temporal sequences:

$$X = \{x_{t-k}, \dots, x_{t-1}, x_t\} \quad \text{where } x_i \in \mathbb{R}^{C \times H \times W}$$

Coupled with spatiotemporal metadata:

$$M = \{m_{t-k}, \dots, m_{t-1}, m_t\} \quad \text{where } m_i = [\text{Lat}, \text{Lon}, \Delta t]$$

2. Roadblocks & Technical Challenges Faced

Throughout the research and engineering lifecycle, five major bottlenecks were encountered:

Roadblock 1: Severe Imbalance in IMD Class Distribution

- **Problem:** Tropical cyclones are rare phenomena; furthermore, high-intensity stages like Extremely Severe Cyclonic Storms (ESCS) and Super Cyclonic Storms (SuCS) account for less than 3% of all recorded satellite frames, whereas Depressions (D) and Deep Depressions (DD) account for over 65%.
- **Symptom:** Standard Cross-Entropy loss caused the network to collapse into predicting the majority classes (D/DD), completely ignoring rapid intensification phases.

Roadblock 2: Timestamp Asynchrony & Ground Truth Label Gaps

- **Problem:** Satellite passes from INSAT-3DR are captured asynchronously every 15–30 minutes, whereas IMD Best Track ground truth observations are logged strictly at 3-hourly or 6-hourly synoptic intervals (00:00, 03:00, 06:00, 12:00, 18:00 UTC).
- **Symptom:** Thousands of raw `.h5` satellite frames had no exact timestamp match in historical cyclone logs. Attempting step-interpolation introduced severe label noise and artificial jumps in wind speed.

Roadblock 3: Sequence Length Variability & Temporal Gaps

- **Problem:** Satellite downlinks frequently suffer packet dropouts, sun-glint blackouts, or calibration cycles, creating irregular time gaps ranging from 15 minutes to over 4 hours between consecutive frames.
- **Symptom:** Fixed recurrent neural networks (RNNs) assume uniform Δt . Feeding variable-interval sequences into an unweighted recurrent unit corrupted temporal feature representations.

Roadblock 4: Overfitting on Landmark Historical Storms

- **Problem:** Training sets heavily featured famous storms (e.g., Fani 2019, Amphan 2020, Tauktae 2021). The model began memorizing specific cloud spiral artifacts unique to Super Cyclone Amphan rather than learning generalizable cyclogenesis physics.
- **Symptom:** High training accuracy (>97%) but poor out-of-sample generalization on newly forming storms in the Arabian Sea.


Roadblock 5: Production Deployment Constraints

- **Problem:** Modern serverless platforms (e.g., Vercel) impose strict 50MB–250MB bundle size limits and execution timeouts (10–60 seconds), prohibiting heavy PyTorch CUDA binaries and multi-gigabyte HDF5 data cubes from running directly in serverless cloud handlers.

3. Engineering Solutions & Methodological Breakthroughs

To address each roadblock rigorously, we introduced the following engineering paradigms:

ARCHITECTURE WORKFLOW



Solution 1: Focal Loss with Class-Frequency Weighting

We replaced standard Cross-Entropy with a generalized Multi-Class Focal Loss formulation with focusing parameter $\gamma = 2.0$ and inverse class-frequency α weights:

$$\mathcal{L}_{\text{Focal}} = -\alpha_t (1 - p_t)^\gamma \log(p_t)$$

This mathematically down-weights easy, well-classified examples (Depressions) and forces model gradient updates to concentrate on hard, rare intense cyclones (VSCS and Super Cyclones).

Solution 2: Cubic Hermite Spline Trajectory Interpolation

To bridge the asynchronous gap between 30-minute INSAT-3DR passes and 6-hourly IMD ground truth records, we implemented Piecewise Cubic Hermite Interpolating Polynomials (PCHIP). Unlike linear interpolation, PCHIP guarantees monotonicity and prevents artificial overshoot in peak sustained wind speeds while preserving true physical rates of central pressure deepening.

Solution 3: Delta-t Aware Temporal Sequence Construction

We developed a dynamic sequence builder that:

- Enforces minimum sequence length $T \geq 2$ and maximum length $T = 10$.
- Discards sequences where the temporal gap between consecutive frames exceeds $\Delta t_{\max} = 90$ minutes.
- Concatenates physical time delta vectors $[\Delta t]$ directly into the GRU input alongside spatial CNN embeddings, allowing the recurrent network to learn velocity derivatives invariant to frame rate.

Solution 4: Stochastic Weight Averaging (SWA) & Heavy Augmentations

To eliminate storm memorization:

- Implemented rotational invariance augmentations (random $\pm 180^\circ$ rotation, horizontal/vertical reflection), reflecting the Coriolis rotation characteristics of the Northern Hemisphere.
- Applied Stochastic Weight Averaging (SWA) over the final 5 training epochs. SWA finds wider, flatter optima in the loss landscape, yielding superior generalization on unseen ocean disturbances.

Solution 5: Decoupled Edge-Cloud Hybrid Architecture

To conquer Vercel's serverless limits:

- Inference Engine:** Runs on dedicated compute (local machine, GPU instances, or Google Colab) executing model inference and writing lightweight JSON telemetry.
- Cloud Source of Truth:** Supabase PostgreSQL acts as the high-speed data broker with sub-millisecond querying and automatic row-level security.
- Frontend Delivery:** Vercel edge nodes serve the Next.js geospatial dashboard with 0ms cold starts and 0MB Python weight overhead.

4. Current Model Architecture & Neural Mechanics

The operational model is implemented in `finetune_cnn_gru_v1_expanded.py` under the class `FixedTemporalChakravatManthan`.

4.1 Spatial CNN Encoder

The spatial encoder processes each individual image frame in the sequence:

- Input Tensor:** $[B \times T \times C \times H \times W]$ where B is batch size, T is sequence length ($2 \leq T \leq 10$), $C = 3$ channels, $H = W = 256$ pixels.
- ConvBlock 1:** 3×32 filters, 3×3 kernel, BatchNorm, ReLU, 2×2 MaxPool.
- ConvBlock 2:** 32×64 filters, 3×3 kernel, BatchNorm, ReLU, 2×2 MaxPool.

- **ConvBlock 3:** 64 to 128 filters, 3 to 3 kernel, BatchNorm, ReLU, 2 to 2 MaxPool.
- **ConvBlock 4:** 128 to 256 filters, 3 to 3 kernel, BatchNorm, ReLU, 2 to 2 MaxPool.
- **Global Pooling & Linear Projection:** AdaptiveAvgPool2d to Linear(256 to 256) to BatchNorm1d to ReLU.
- **Spatial Feature Dimension:** $d_{\text{cnn}} = 256$.

4.2 Spatio-Temporal Recurrent Core (GRU)

- **Recurrent Layer:** 2-Layer Gated Recurrent Unit (GRU).
- **Input Dimension:** $d_{\text{in}} = d_{\text{cnn}} + d_{\text{meta}} = 256 + 3 = 259$ (incorporating normalized latitude, longitude, and elapsed time).
- **Hidden State Dimension:** $d_{\text{hidden}} = 128$.
- **Recurrent Dropout:** $p = 0.3$.
- The final hidden state $h_T \in \mathbb{R}^{B \times 128}$ encapsulates the historical trajectory, cloud spiral growth rate, and thermodynamic momentum of the cyclone.

4.3 Multi-Task Output Heads

1. Intensity Classification Head:

- Architecture: Linear(128 to 64) to ReLU to Dropout(0.4) to Linear(64 to 8).
- Classes: 8 official IMD categories:

$\text{IMD Stages} = \{\text{TD}, \text{D}, \text{DD}, \text{CS}, \text{SCS}, \text{VSCS}, \text{ESCS}, \text{SuCS}\}$

2. Neural Wind Regression Head:

- Architecture: Linear(128 to 32) to ReLU to Linear(32 to 1).
- Predicts continuous Maximum Sustained Wind Speed V_{max} in knots.

3. Physical Atmospheric Pressure Derivation:

- Rather than treating pressure as an independent unconstrained variable, the model couples V_{max} through the Atkinson-Holliday empirical cyclostrophic formulation calibrated for the Bay of Bengal:

$$P_{\text{min}} = 1010 - \left(\frac{V_{\text{max}}}{2.3} \right)^{1.33} \quad [\text{hPa}]$$


4. Spatial Geometry Estimator:

- Predicts Outer Gale Wind Radius (R_{34} / Swath), Central Dense Overcast (CDO) core radius, and Eye Wall diameter.

5. Preprocessing & Radiometric Calibration Pipeline

The satellite data ingestion pipeline handles raw ISRO MOSDAC INSAT-3DR L1C HDF5 archives:

ARCHITECTURE WORKFLOW



1. **Spectral Band Selection:** Ingests the Thermal Infrared-1 (TIR-1, 10.8\,\mu\text{m}) channel. This channel is critical because it operates continuously 24 hours a day (day and night) and directly correlates cloud-top temperatures with storm convection height.
2. **Brightness Temperature Conversion:** Digital numbers (DN) are converted to top-of-atmosphere radiance and subsequently to equivalent blackbody brightness temperatures in Kelvin using the Planck inversion function.
3. **Spatial Normalization:** Images are cropped and resized to 256 \times 256 pixels using bilinear interpolation, with zero-padding applied if edge coordinates intersect observation limits.

6. Model Verification & Benchmark Results

The final trained checkpoint (chakravat_manthan_cnn_gru_fixed_best.pt) achieved benchmark-grade results on the independent test split:

Evaluation Metric	Score	Operational Significance
Exact Category Accuracy	91.94%	Correctly identifies exact IMD stage out of 8 classes.
Adjacent Category Accuracy (± 1 stage)	98.90%	Crucial for operational safety: errors never jump more than one category.
Mean Absolute Error (MAE) - Wind Speed	4.2 knots	Well within IMD's operational tolerance threshold (± 5 kt).
Root Mean Square Error (RMSE) - Pressure	3.8 hPa	Captures rapid pressure drops during intensification phases.
Inference Latency (GPU - T4 / RTX)	18 ms	Real-time throughput capability.
Inference Latency (CPU - Intel/AMD)	64 ms	Fully deployable on lightweight edge CPU servers without dedicated GPU.

7. Current Operational Condition & Live Status

As of **September 28, 2026**:

1. **Model Weights:** Safely pushed and hosted in the GitHub repository (``backend/checkpoints/``).
2. **Decoupled Two-Database Zero-Loss Buffer Pipeline:**
 - **Database 1 (Buffer Queue):** Dedicated database (``chakravat-ingestion-buffer``) ingesting raw ISRO MOSDAC passes and storing 512x512 crops with status flags (``PENDING`` -> ``PROCESSED``). Guarantees zero dropped frames during shared CI runner delays or network congestion.
 - **Database 2 (Live Telemetry):** Dedicated operational database (``cyclone_live``) serving real-time inferences to the Next.js Vercel frontend.
3. **Pan-Asia Meteorological Coverage:** Geospatial viewing domain expanded across the entire Asian continent and oceanic basins (40.0°E to 145.0°E , -10.0°S to 48.0°N), from the Arabian Peninsula to Japan, China, Southeast Asia, and the Western Pacific.
4. **Physical Wind Streamlines:** Real atmospheric GFS/ECMWF streamlines dynamically calculated across the entire Pan-Asia grid; artificial vortex rotation is triggered strictly when an actual tropical storm ($\geq 28\text{ kt}$) is verified.
5. **Live Production Dashboard:** Hosted on Vercel at <https://chakravat-manthan-live.vercel.app> with zero template branding and private admin telemetry monitoring.

8. Conclusion & Future Roadmap

The Chakravat Manthan machine learning framework represents a robust, meteorologically sound, and production-tested system for tropical cyclone monitoring. By combining physical atmospheric constraints with spatio-temporal neural dynamics and a decoupled zero-loss buffer pipeline, the platform provides early warning capabilities that bridge the gap between complex satellite telemetry and life-saving disaster management decisions.

Planned Next Phases:

- Integration of Scatterometer (SCATSAT-1 / EOS-06) ocean surface roughness measurements.
- Physics-Informed Neural Networks (PINNs) incorporating the Navier-Stokes vorticity transport equation directly into the loss function.
- Automated 24/7 cloud cron execution via GitHub Actions for continuous MOSDAC pass processing.